

# Performance benchmarks for common circuit simulators

Evan

August 4, 2019

## **Abstract**

abstract here

# Contents

<b>1</b>	<b>Methods</b>	<b>3</b>
<b>2</b>	<b>Simulator methods</b>	<b>3</b>
2.1	General unitary matrix multiplication . . . . .	3
2.2	Clifford circuits . . . . .	4
2.2.1	Example: $X_k$ . . . . .	4
2.2.2	General Clifford operations . . . . .	5
<b>3</b>	<b>Benchmark methods</b>	<b>5</b>
<b>4</b>	<b>Results</b>	<b>5</b>
<b>A</b>	<b>Appendix 1: Einsum operations</b>	<b>5</b>

# 1 Methods

## 2 Simulator methods

### 2.1 General unitary matrix multiplication

A complete wavefunction simulator (cite Biamonte for what that means; I'll be reffing tensor networks anyways) generally applies the action of a gate as an einsum to the qubit axes acted on by the gate.

Worked example: Computing  $U_2 U_1 \vec{\psi}$  For example, let  $U_1$  be a unitary over a subset of  $k_1$  and  $U_2$  be qubits on a state defined over  $k_2$ , and the target state  $|\psi\rangle$  be defined over  $n$  qubits. Then we have:

$$\vec{\psi} \in \mathbb{C}^{2^n} \tag{1}$$

$$U_1 \in \mathbb{C}^{2^{k_1} \times 2^{k_1}} \tag{2}$$

$$U_2 \in \mathbb{C}^{2^{k_2} \times 2^{k_2}} \tag{3}$$

$$\tag{4}$$

The following list details approaches to computing  $U_2 U_1 |\psi\rangle$  in order of worst to best:

1. **kronecker product + matmul:** This requires resizing the matrices to  $n$  dimensions then performing matrix multiplication

- (a) kronecker product  $U_1 \rightarrow U'_1 \in \mathbb{C}^{2^n \times 2^n}$  using  $I_2$  ( $\mathbb{O}(2_1^k 2_1^{2n-k_1} 2^{n-k_1}) = \mathbb{O}(2^{2n})$ )

- (b) kronecker product  $U_2 \rightarrow U'_2 \in \mathbb{C}^{2^n \times 2^n}$  using  $I_2$  ( $\mathbb{O}(2_2^k 2_2^{2n-k_2} 2^{n-k_2}) = \mathbb{O}(2^{2n})$ )

- (c) matmul  $U'_2 U'_1 \rightarrow U_f \in \mathbb{C}^{2^n \times 2^n}$  ( $\mathbb{O}(2^{3n})$ )<sup>1</sup>

- (d) matmul  $U_f \vec{\psi}$  ( $\mathbb{O}(2^{2n})$ )

2. **kronecker product + matmul, associativity:** This swaps (c) and (d), taking advantage of the fact that  $U \vec{\psi}$  has complexity  $\mathbb{O}(2^{2n})$  for  $U \in \mathbb{C}^{2^n \times 2^n}$ , meaning both products cost  $\mathbb{O}(2^{2n+1})$  instead of  $\mathbb{O}(2^{3n} + 2^{2n})$ . Conceptually this results from not having to calculate the full intermediate  $U_f$ .

3. **einsum:** An einstein summation allows general tensor transformations using repeated (summation) and permuted (free) indices (See Appendix A

---

<sup>1</sup>Assuming “schoolbook” matrix multiplication algorithm. For square matrix multiplication this can be improved to  $\mathbb{O}(2^{2.807n})$  or even  $\mathbb{O}(2^{2.37n})$  with ML but that's not the point of this worst-case example

for some common examples). The austere implementation results in a complexity of<sup>2</sup>:

$$\mathbb{O} \left( \left( \prod_i^{N_{\text{free}}} d_i \right) \left( \prod_i^{N_{\text{sum}}} d_i \right) \right) \quad (5)$$

where  $N_{\text{free}}$  is the number of free indices (indices of output),  $N_{\text{sum}}$  is the number of unique input indices between input tensors, and  $d_i$  is the size of the axis corresponding to the  $i$ -th index.<sup>2</sup>

Let unitaries and wavefunctions be represented as tensors of the shape:

$$U \in \mathbb{C}^{\overbrace{2 \times 2 \times \dots}^{k \text{ times}}} \quad (6)$$

$$\vec{\psi} \in \mathbb{C}^{\overbrace{2 \times \dots}^{n \text{ times}}} \quad (7)$$

Then the einsum computing  $U\vec{\psi}$  is indexed as

$$i_1 \dots i_k, i_1 \dots i_n \rightarrow i_1 \dots i_n$$

has  $k$  summation indices and  $(n - k)$  free indices, resulting in a complexity of  $\mathbb{O}(2^n)$  since every tensor axis is dimension two. This is a two-fold speedup over square matrix math. Conceptually this results from  $?!?!?!?!.$

The complexity of method (3) is obviously a lower bound since *computing* all of the elements in  $\vec{\psi}' = U\vec{\psi}$  elements can require no less than  $2^n$  computations. However this is cannot be the minimum complexity of processing (FIXME: computing vs processing vs solving..?!?!?!?)  $U\vec{\psi}$ , since the simple counterexample of  $I\vec{\psi}$  requires exactly 0 computations. This motivates even more streamlined unitary action, which will be introduced in the following sections.

## 2.2 Clifford circuits

### 2.2.1 Example: $X_k$

This section begins with an example of applying a permutation to a tensor subspace (taken from the `apply_unitary` from the `Cirq` library). Let  $X_k$  be the Pauli-X gate acting on the qubit indexed “ $k$ ”. Define the operation `slice_for_bitstring( $T, \{i \text{ for } i=1\dots k\}, \vec{s}$ )` to access elements of tensor  $T$  for which the subspace corresponding to qubits  $\{0, 1, \dots, k\}$  is represented by the (little-endian) bitstring  $\vec{s}$  (and so  $|\vec{s}| = |\{i_k\}|$ ). TODO: example of this....

This slices for a subset of  $2^{n-k}$  elements, but for an array stored with a known memory layout and stride the slice can be accessed in  $\mathbb{O}(k)$  time (DEMONSTRATED EMPIRICALLY, PROVE THIS?). Then the action of  $X_k$  on  $\vec{\psi}$  can be computed in two steps using a memory buffer  $T'$  sized the same as  $T$ :

<sup>2</sup>For example, the einsum string `ik,kj->ij` multiplies two matrices using free indices  $\mathbf{i}, \mathbf{j}$  and summation index  $\mathbf{k}$ ; if  $\mathbf{i}, \mathbf{j}, \mathbf{k}$  index axes of sizes  $\ell, m, n$  respectively, the complexity of this einsum is  $\mathbb{O}((d_i d_j)(d_k)) = \mathbb{O}(\ell m n)$ , the expected complexity of matrix multiplication.

- Compute  $S_0 = \text{slice\_for\_bitstring}(T, \{k\}, (0))$  and  $S_1 = \text{slice\_for\_bitstring}(T, \{k\}, (0))$  which slice  $T$  for all elements in which the  $k$ -th qubit is a “0” or “1” respectively
- Set  $T'[S_1] = T[S_0]$  and  $T'[S_0] = T[S_1]$ . In plain English, this sets the amplitude for each basis state in  $T'$  to be the same as the amplitude in  $T$  where the  $k$ -th qubit of that basis state is flipped.

### 2.2.2 General Clifford operations

Section 2.2.1 demonstrated how a single-qubit gate acting on an  $n$ -qubit state can be implemented in  $\mathcal{O}(n)$ , which is exponential improvement over the best matrix-style operations (at the expense of exponential memory overhead in the number of qubits, i.e. the buffer state).

- Clifford circuits are easy to simulate (Gottesman). Computationally these are relatively sparse matrices.

Here we review an algorithmic app

## 3 Benchmark methods

This repo uses the PYTEST-BENCHMARK fixture to profile python code.

## 4 Results

### Appendix A Appendix 1: Einsum operations

Some common einsum-compatible operations: transpose, inner product, matrix multiplication, trace.