

Performance benchmarks for common circuit simulators

Evan

August 7, 2019

Abstract

abstract here

Contents

1	Methods	3
2	Simulator methods	3
2.1	General unitary matrix multiplication	3
2.2	Nonlinear mappings	4
2.2.1	Example: X_k	4
2.2.2	General Clifford operations	5
3	Benchmark methods	5
4	Results	5
A	Appendix 1: Einsum operations	5

1 Methods

2 Simulator methods

2.1 General unitary matrix multiplication

A complete wavefunction simulator (cite Biamonte for what that means; I'll be reffing tensor networks anyways) generally applies the action of a gate as an einsum to the qubit axes acted on by the gate.

Worked example: Computing $U_2 U_1 \vec{\psi}$ For example, let U_1 be a unitary over a subset of k_1 and U_2 be qubits on a state defined over k_2 , and the target state $|\psi\rangle$ be defined over n qubits. Then we have:

$$\vec{\psi} \in \mathbb{C}^{2^n} \tag{1}$$

$$U_1 \in \mathbb{C}^{2^{k_1} \times 2^{k_1}} \tag{2}$$

$$U_2 \in \mathbb{C}^{2^{k_2} \times 2^{k_2}} \tag{3}$$

$$\tag{4}$$

The following list details approaches to computing $U_2 U_1 |\psi\rangle$ in order of worst to best:

1. **kroncker product + matmul:** This requires resizing the matrices to n dimensions then performing matrix multiplication

- (a) kroncker product $U_1 \rightarrow U'_1 \in \mathbb{C}^{2^n \times 2^n}$ using I_2 ($\mathbb{O}(2_1^k 2_1^{2n-k_1} 2^{n-k_1}) = \mathbb{O}(2^{2n})$)

- (b) kroncker product $U_2 \rightarrow U'_2 \in \mathbb{C}^{2^n \times 2^n}$ using I_2 ($\mathbb{O}(2_2^k 2_2^{2n-k_2} 2^{n-k_2}) = \mathbb{O}(2^{2n})$)

- (c) matmul $U'_2 U'_1 \rightarrow U_f \in \mathbb{C}^{2^n \times 2^n}$ ($\mathbb{O}(2^{3n})$)¹

- (d) matmul $U_f \vec{\psi}$ ($\mathbb{O}(2^{2n})$)

2. **kroncker product + matmul, associativity:** This swaps (c) and (d), taking advantage of the fact that $U \vec{\psi}$ has complexity $\mathbb{O}(2^{2n})$ for $U \in \mathbb{C}^{2^n \times 2^n}$, meaning both products cost $\mathbb{O}(2^{2n+1})$ instead of $\mathbb{O}(2^{3n} + 2^{2n})$. Conceptually this results from not having to calculate the full intermediate U_f .

3. **einsum:** An einstein summation allows general tensor transformations using repeated (summation) and permuted (free) indices (See Appendix A

¹Assuming "schoolbook" matrix multiplication algorithm. For square matrix multiplication this can be improved to $\mathbb{O}(2^{2.807n})$ or even $\mathbb{O}(2^{2.37n})$ with ML but that's not the point of this worst-case example

for some common examples). The austere implementation results in a complexity of²:

$$\mathbb{O} \left(\left(\prod_i^{N_{\text{free}}} d_i \right) \left(\prod_i^{N_{\text{sum}}} d_i \right) \right) \quad (5)$$

where N_{free} is the number of free indices (indices of output), N_{sum} is the number of unique input indices between input tensors, and d_i is the size of the axis corresponding to the i -th index.²

Let unitaries and wavefunctions be represented as tensors of the shape:

$$U \in \mathbb{C}^{\overbrace{2 \times 2 \times \dots}^{k \text{ times}}} \quad (6)$$

$$\vec{\psi} \in \mathbb{C}^{\overbrace{2 \times \dots}^{n \text{ times}}} \quad (7)$$

with the restriction that $n \geq 2k$ to ensure compatible the state and matrix are compatible for multiplication. Then the einsum computing $U\vec{\psi}$ is indexed as³

$$i_1 \dots i_k, i'_1 \dots i'_k i_1 \dots i_n \rightarrow i_{k+1} \dots i_n i'_1 \dots i'_k$$

has k summation indices and $(n-2k)$ free indices, resulting in a complexity of $\mathbb{O}(2^{n+k})$ since every tensor axis is dimension two.

The complexity of method (3) does not saturate the lower bound for computing all of the elements in $\vec{\psi}' = U\vec{\psi}$ elements, which can require as few as 2^n operations (namely, $I\vec{\psi}$ using sparse multiplication saturates this bound). This motivates even more streamlined unitary action, which will be introduced in the following sections.

2.2 Nonlinear mappings

2.2.1 Example: X_k

This section begins with an example of applying a permutation to a tensor subspace (taken from the `apply_unitary` from the `CIRQ` library). Let X_k

²For example, the einsum string `ik,kj->ij` multiplies two matrices using free indices **i**, **j** and summation index **k**; if **i**, **j**, **k** index axes of sizes ℓ, m, n respectively, the complexity of this einsum is $\mathbb{O}((d_i d_j)(d_k)) = \mathbb{O}(\ell m n)$, the expected complexity of matrix multiplication.

³Three worked examples:

- (a) single-qubit matrix M acting on two-qubit state ψ ($k = 1, n = 2$): $(M\psi)_{jk} = M_{ij}\psi_{ik}$ has one summation index **i** and two free indices **j,k**, for $\mathbb{O}(2^3)$
- (b) two-qubit matrix M acting on two-qubit state ψ ($k = 2, n = 2$): $(M\psi)_{\ell m} = M_{jk\ell m}\psi_{jk}$ has two summation indices **j,k** and two free indices ℓ, m , for $\mathbb{O}(2^4)$
- (c) two-qubit matrix M acting on three-qubit state ψ ($k = 2, n = 3$): $(M\psi)_{\ell mn} = M_{jk\ell m}\psi_{jkn}$ has two summation indices **j,k** and three free indices ℓ, m, n , for $\mathbb{O}(2^5)$

be the Pauli-X gate acting on the qubit indexed “k”. Define the operation `slice_for_bitstring(T, {i for i=1...k}, \vec{s})` to access elements of tensor T for which the subspace corresponding to qubits $\{0, 1, \dots, k\}$ is represented by the (little-endian) bitstring \vec{s} (and so $|\vec{s}| = |\{i_k\}|$). TODO: example of this....

This slices for a subset of 2^{n-k} elements. Then the action of X_k on $\vec{\psi}$ can be computed in two steps using a memory buffer T' sized the same as T (namely, 2^n):

- Compute $S_0 = \text{slice_for_bitstring}(T, \{k\}, (0))$ and $S_1 = \text{slice_for_bitstring}(T, \{k\}, (1))$ which slice T for all elements in which the k-th qubit is a “0” or “1” respectively
- Set $T'[S_1] = T[S_0]$ and $T'[S_0] = T[S_1]$ at a cost of $\mathcal{O}(2^{n-1})$ each, for a total complexity requirement of $\mathcal{O}(2^n)$. In plain English, this sets the amplitude for each basis state in T' to be the same as the amplitude in T where the k-th qubit of that basis state is flipped.

Section 2.2.1 demonstrated how a certain single-qubit gate acting on an n-qubit state can be implemented in $\mathcal{O}(2^n)$, which is at least $2\times$ improvement over the best matrix-style operations (at the expense of exponential memory overhead in the number of qubits, i.e. the uninitialized buffer state).

TODO

2.2.2 General Clifford operations

- Clifford circuits are easy to simulate (Gottesman). Computationally these are relatively sparse matrices.

Here we review an algorithmic app

3 Benchmark methods

This repo uses the PYTEST-BENCHMARK fixture to profile python code.

4 Results

Appendix A Appendix 1: Einsum operations

Some common einsum-compatible operations: transpose, inner produc, matrix multiplication, trace.